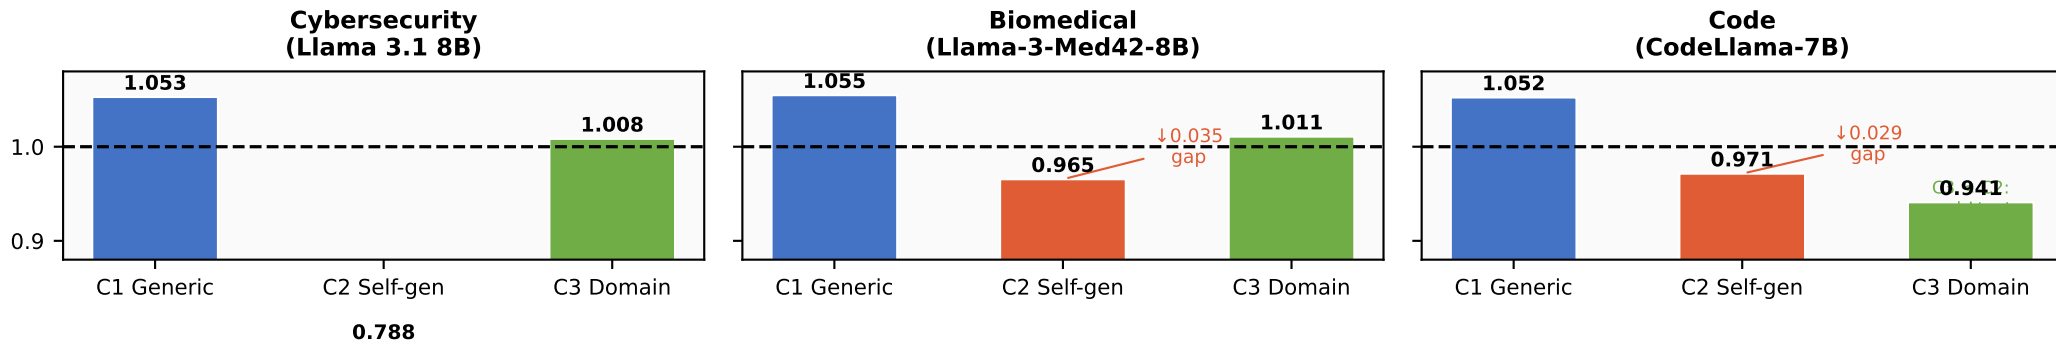


Fig. 5: Per-channel input activation range under each calibration condition relative to MMLU inference. C2 underestimates in all domains; C3 $\approx$ MMLU for cyber/med, but also underestimates for code (code text is too syntactically narrow to cover MMLU range).

Mean per-channel range ratio
vs. MMLU inference



■ C1 Generic ■ C2 Self-gen ■ C3 Domain --- MMLU inference (ratio = 1.0)