

Fig. 5: Per-channel input activation range under each calibration condition relative to MMLU inference. C2 underestimates in all domains; C3 $\approx$ MMLU for cyber/med, but also underestimates for code (code text is too syntactically narrow to cover MMLU range).

